

```
In [599... # Kaggle Competitions: Outflow of user (DeepLearningSchool)
# https://www.kaggle.com/competitions/advanced-dls-spring-2021
```

```
In [ ]: # Install Library
!pip install catboost
!pip install xgboost
```

```
In [601... # imports
import pandas as pd
import numpy as np
import re

import seaborn as sns
from matplotlib import pyplot as plt
%matplotlib inline
```

Load data

```
In [602... # pathes/filenames
PATH_PROJECT = ""
PATH_DATASET = PATH_PROJECT + "dataset/"

train_filename = PATH_DATASET + "train.csv"
test_filename = PATH_DATASET + "test.csv"
submission_filename = PATH_DATASET + "test_submission.csv"
```

```
In [603... # Load dataset
outflow_users_df = pd.read_csv(train_filename)
test_outflow_users_df = pd.read_csv(test_filename)

# Numeric columns
num_cols = ['ClientPeriod', 'MonthlySpending', 'TotalSpent']
# Categorical columns
cat_cols = ['Sex', 'IsSeniorCitizen', 'HasPartner', 'HasChild',
            'HasPhoneService',
            'HasMultiplePhoneNumbers',
            'HasInternetService',
            'HasOnlineSecurityService',
            'HasOnlineBackup', 'HasDeviceProtection', 'HasTechSupportAccess', 'HasOr',
            'HasMovieSubscription', 'HasContractPhone', 'IsBillingPaperless', 'Paymer
            ]
# Target column
target_colname = "Churn"

# show shape
print(outflow_users_df.shape, test_outflow_users_df.shape)
outflow_users_df.head(3)

(5282, 20) (1761, 19)
```

Out[603]:

	ClientPeriod	MonthlySpending	TotalSpent	Sex	IsSeniorCitizen	HasPartner	HasChild	HasPartner
0	55	19.50	1026.35	Male	0	Yes	Yes	
1	72	25.85	1872.2	Male	0	Yes	No	
2	1	75.90	75.9	Male	0	No	No	

Data preprocessing

```
In [604... # Change synonyms test/train datasets
def remove_synonyms(df):
    # remove synonyms phone service
    list_phone_service = ["HasMultiplePhoneNumbers"]
    for phone_service in list_phone_service:
        df[phone_service] = np.where(df[phone_service] == "No phone service", "No",

    # remove synonyms internet service
    list_internet_service = ["HasOnlineSecurityService", "HasOnlineBackup", "HasDev
        "HasTechSupportAccess", "HasOnlineTV", "HasMovieSubscr
    ]
    for internet_service in list_internet_service:
        df[internet_service] = np.where(df[internet_service] == "No internet service", "No",

    return df

outflow_users_df = remove_synonyms(outflow_users_df)
test_outflow_users_df = remove_synonyms(test_outflow_users_df)
```

```
In [605... # Correct Column "TotalSpent" contains spaces (" ") and ClientPeriod with zeros value
def correct_total_spent(df):
    total_spent_column = "TotalSpent"
    period_column = "ClientPeriod"
    monthly_spending_column = "MonthlySpending"

    # replace empty string on zero value (0)
    df[total_spent_column] = [float(re.sub(r"\s+", "0", x)) for x in df[total_spent_column]]

    # replace zero value on calc value
    df[total_spent_column] = np.where(df[total_spent_column] == 0, df[period_column] * df[monthly_spending_column], df[total_spent_column])

    # if exist zero -> replace zero value on mean value
    df[total_spent_column] = np.where(df[total_spent_column] == 0, df[total_spent_column].mean(), df[total_spent_column])
    # if "ClientPeriod" exist zero -> replace zero value on TotalSpent // MonthlySpending
    df[period_column] = np.where(df[period_column] == 0, df[total_spent_column] // df[monthly_spending_column], df[period_column])
    return df

outflow_users_df = correct_total_spent(outflow_users_df)
test_outflow_users_df = correct_total_spent(test_outflow_users_df)
```

```
In [606... outflow_users_df.info()
outflow_users_df.describe()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 5282 entries, 0 to 5281
Data columns (total 20 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   ClientPeriod                          5282 non-null   float64
1   MonthlySpending                       5282 non-null   float64
2   TotalSpent                            5282 non-null   float64
3   Sex                                    5282 non-null   object
4   IsSeniorCitizen                       5282 non-null   int64
5   HasPartner                            5282 non-null   object
6   HasChild                              5282 non-null   object
7   HasPhoneService                       5282 non-null   object
8   HasMultiplePhoneNumbers               5282 non-null   object
9   HasInternetService                   5282 non-null   object
10  HasOnlineSecurityService              5282 non-null   object
11  HasOnlineBackup                       5282 non-null   object
12  HasDeviceProtection                   5282 non-null   object
13  HasTechSupportAccess                  5282 non-null   object
14  HasOnlineTV                           5282 non-null   object
15  HasMovieSubscription                  5282 non-null   object
16  HasContractPhone                      5282 non-null   object
17  IsBillingPaperless                    5282 non-null   object
18  PaymentMethod                         5282 non-null   object
19  Churn                                 5282 non-null   int64
dtypes: float64(3), int64(2), object(15)
memory usage: 825.4+ KB
```

Out[606]:

	ClientPeriod	MonthlySpending	TotalSpent	IsSeniorCitizen	Chum
count	5282.000000	5282.000000	5282.000000	5282.000000	5282.000000
mean	32.507952	64.924754	2291.383808	0.159409	0.262022
std	24.591755	30.176464	2267.079962	0.366092	0.439776
min	1.000000	18.250000	18.800000	0.000000	0.000000
25%	9.000000	35.462500	401.587500	0.000000	0.000000
50%	29.000000	70.400000	1413.650000	0.000000	0.000000
75%	55.000000	90.050000	3795.212500	0.000000	1.000000
max	116.000000	118.750000	8684.800000	1.000000	1.000000

In [607...

```
# Check is NaN values
outflow_users_df.isna().sum(), test_outflow_users_df.isna().sum()
```

```

Out[607]: (ClientPeriod      0
MonthlySpending      0
TotalSpent           0
Sex                  0
IsSeniorCitizen      0
HasPartner           0
HasChild             0
HasPhoneService      0
HasMultiplePhoneNumbers 0
HasInternetService   0
HasOnlineSecurityService 0
HasOnlineBackup      0
HasDeviceProtection  0
HasTechSupportAccess 0
HasOnlineTV          0
HasMovieSubscription 0
HasContractPhone     0
IsBillingPaperless   0
PaymentMethod        0
Churn                0
dtype: int64,
ClientPeriod      0
MonthlySpending      0
TotalSpent           0
Sex                  0
IsSeniorCitizen      0
HasPartner           0
HasChild             0
HasPhoneService      0
HasMultiplePhoneNumbers 0
HasInternetService   0
HasOnlineSecurityService 0
HasOnlineBackup      0
HasDeviceProtection  0
HasTechSupportAccess 0
HasOnlineTV          0
HasMovieSubscription 0
HasContractPhone     0
IsBillingPaperless   0
PaymentMethod        0
dtype: int64)

```

```

In [608... # Show columns name
cat_cols, num_cols

```

```
Out[608]: ([ 'Sex',
             'IsSeniorCitizen',
             'HasPartner',
             'HasChild',
             'HasPhoneService',
             'HasMultiplePhoneNumbers',
             'HasInternetService',
             'HasOnlineSecurityService',
             'HasOnlineBackup',
             'HasDeviceProtection',
             'HasTechSupportAccess',
             'HasOnlineTV',
             'HasMovieSubscription',
             'HasContractPhone',
             'IsBillingPaperless',
             'PaymentMethod'],
            ['ClientPeriod', 'MonthlySpending', 'TotalSpent'])
```

One-Hot Encoder Categorical columns

```
In [609... # One-Hot Encoder ("OneHotEncoder" better than "get_dummies")
from sklearn.preprocessing import OneHotEncoder

sk_ohe = OneHotEncoder(handle_unknown="ignore", drop="first", sparse=False, dtype=int)
sk_ohe = sk_ohe.fit(outflow_users_df[cat_cols])
df_cat_ohe_sklearn = pd.DataFrame(sk_ohe.transform(outflow_users_df[cat_cols]), columns=sk_ohe.get_feature_names_out())
df_cat_ohe_sklearn_testfile = pd.DataFrame(sk_ohe.transform(test_outflow_users_df[cat_cols]), columns=sk_ohe.get_feature_names_out())

# show shape
print(df_cat_ohe_sklearn.shape)
df_cat_ohe_sklearn.head(3)
```

(5282, 20)

```
Out[609]:
```

	Sex_Male	IsSeniorCitizen_1	HasPartner_Yes	HasChild_Yes	HasPhoneService_Yes	HasMultiplePh
0	1	0	1	1	1	
1	1	0	1	0	1	
2	1	0	0	0	1	


Data analysis

```
In [610... df_analysis = pd.concat([outflow_users_df[num_cols], df_cat_ohe_sklearn, y_data], axis=1)
```

Count values

Target disbalanced!

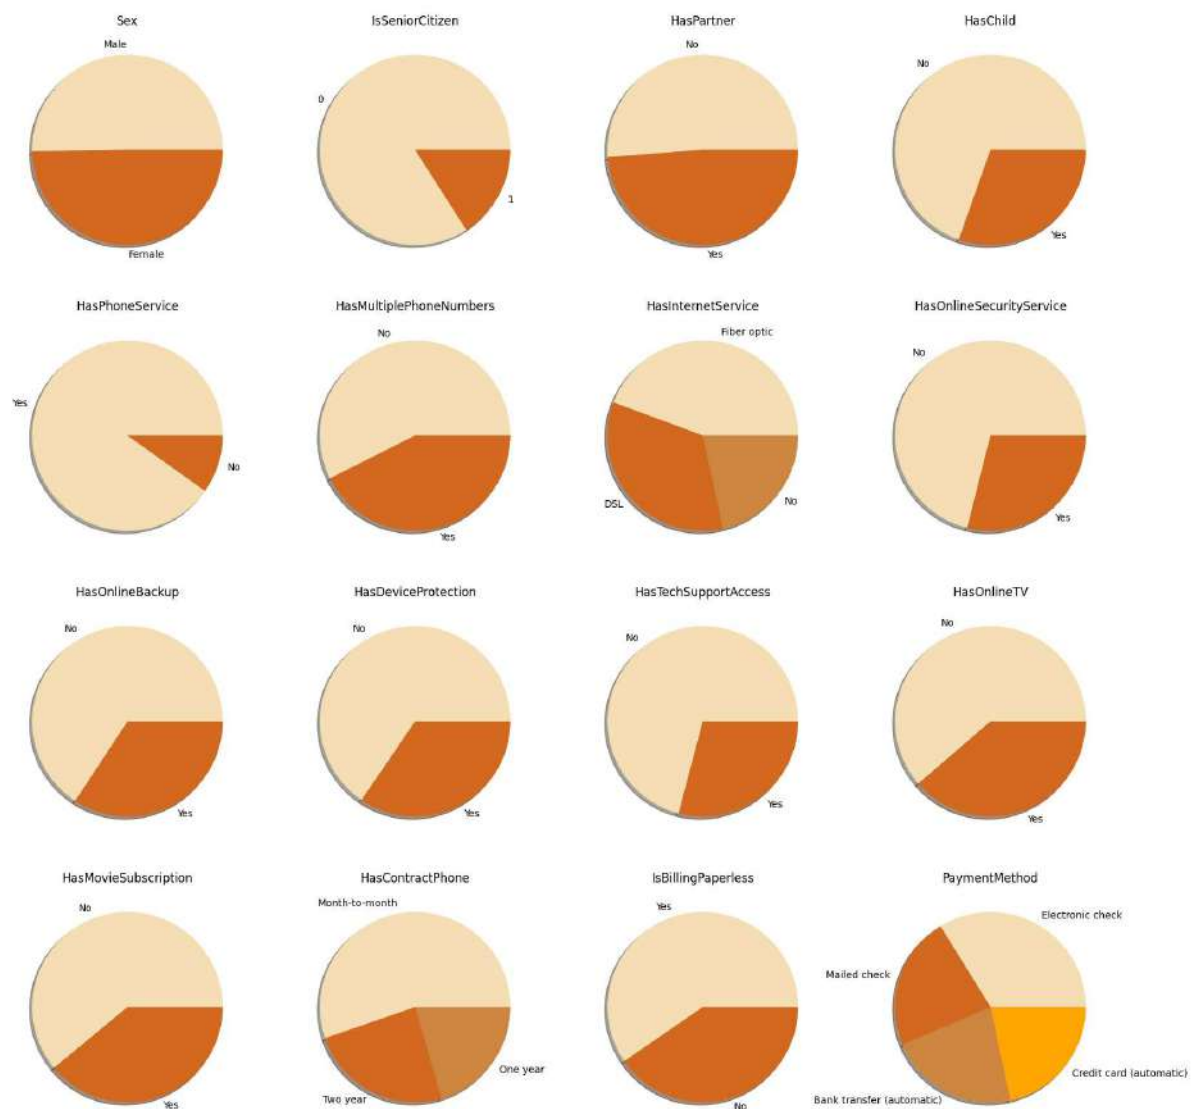
```
In [611... # Show target counts
plt.figure(figsize=(2,2))
plt.pie(y_data.value_counts(), labels=[0, 1], colors=['wheat', 'chocolate'], explode=[0.1, 0.1])
plt.show()
```



In [612...

```
# Show count values for categorical
fig, axes = plt.subplots(4, 4, figsize=(20, 20))
ax = axes.ravel()
fig.suptitle('Number of values for each category', fontsize=16)
for i, cat in enumerate(cat_cols):
    x = outflow_users_df[cat].value_counts()
    index = x.index
    ax[i].pie(labels=index, x=x, colors=['wheat', 'chocolate', 'peru', 'orange'], s
    ax[i].set_title(cat)
```

Number of values for each category



Histogram analysis


Result:

1. The longer the ClientPeriod, the less the Churn. The largest churn of customers is observed in the first months
2. Most regular customers enjoy inexpensive tariffs.
3. Regular customers bring more profit.


```
In [613... # Show histplot for numeric columns
for column in num_cols:
    print(f"Column name: {column}")
    sns.histplot(data=df_analysis, x=column, hue=target_colname, palette="pastel",
    plt.show()

# Show boxplot for numeric columns
for column in num_cols:
    print(f"Column name: {column}")
    sns.boxplot(df_analysis[column], palette='pastel')
    plt.show()
```


Column name: ClientPeriod




Column name: MonthlySpending




Column name: TotalSpent




Column name: ClientPeriod



Column name: MonthlySpending



Column name: TotalSpent



Anomaly detection


```
In [ ]: # Correct anomaly data
upper_limit_ClientPeriod = 75
df_analysis['ClientPeriod'] = np.where(df_analysis['ClientPeriod'] > upper_limit_C
```

Correlation matrix

```
In [615... # correlation = pd.concat([x_data,y_data], axis=1).corr()
correlation = df_analysis.corr()

plt.figure(figsize=(10, 8))
sns.heatmap(
    correlation,
    xticklabels=correlation.columns,
    yticklabels=correlation.columns,
    cmap="RdBu",
    center=0,
)

plt.title("Correlation matrix", fontsize=22)
plt.show()
```



Join preproc columns

```
In [616... # Concatenate OHE columns + numeric columns
x_data = pd.concat([outflow_users_df[num_cols], df_cat_ohe_sklearn],axis=1)
x_data_testfile = pd.concat([test_outflow_users_df[num_cols], df_cat_ohe_sklearn_te

y_data = outflow_users_df[target_colname]

# show shape
print(x_data.shape, x_data_testfile.shape, y_data.shape)
x_data.head(3)

(5282, 23) (1761, 23) (5282,)
```

Out[616]:

	ClientPeriod	MonthlySpending	TotalSpent	Sex_Male	IsSeniorCitizen_1	HasPartner_Yes	HasCh
0	55.0	19.50	1026.35	1	0	1	
1	72.0	25.85	1872.20	1	0	1	
2	1.0	75.90	75.90	1	0	0	

3 rows × 23 columns

Train/Test split

```
In [617... # Train/Test split
from sklearn.model_selection import train_test_split
x_train, x_test, y_train, y_test = train_test_split(x_data, y_data, test_size=0.1,

# show shape
print(x_data.shape, y_data.shape, x_train.shape, x_test.shape, y_train.shape, y_test.shape)
x_train.head(3)

(5282, 23) (5282,) (4753, 23) (529, 23) (4753,) (529,)
```

```
Out[617]:
```

	ClientPeriod	MonthlySpending	TotalSpent	Sex_Male	IsSeniorCitizen_1	HasPartner_Yes	HasPartner_No
1106	2.0	76.5	162.45	1	0	0	1
2095	27.0	79.5	2180.55	1	0	0	1
4469	72.0	104.8	7470.10	1	0	1	0

3 rows × 23 columns

Normalization Numeric columns

```
In [618... from sklearn.preprocessing import StandardScaler
if len(num_cols) > 0:
    scaler = StandardScaler()
    # Fit on train data!
    scaler.fit(x_train[num_cols])
    x_train[num_cols] = scaler.transform(x_train[num_cols])
    x_test[num_cols] = scaler.transform(x_test[num_cols])

    scaler_all = StandardScaler()
    scaler_all.fit(x_data[num_cols])
    x_data_normalize = scaler_all.transform(x_data[num_cols])
    x_data_testfile_normalize = scaler_all.transform(x_data_testfile[num_cols])

# show shape
print(x_data.shape, y_data.shape, x_train.shape, x_test.shape, y_train.shape, y_test.shape)

(5282, 23) (5282,) (4753, 23) (529, 23) (4753,) (529,)
```

Create models

```
In [619... # Load models
from sklearn.linear_model import LogisticRegression
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import roc_auc_score
from sklearn.model_selection import cross_val_predict
import catboost
```

Search feature importances


```
In [620... rf_importances = RandomForestClassifier(max_depth=7, n_estimators=200, n_jobs=-1)
rf_importances.fit(x_train, y_train)

predict_rf_importances = rf_importances.predict_proba(x_test)[: , 1]

print(roc_auc_score(y_test, predict_rf_importances))
feat_importances = pd.Series(rf_importances.feature_importances_, index=x_train.columns)
```

```
feat_importances.nlargest(x_train.shape[0]).plot(kind='barh', title="Feature importances")
plt.show()
```

0.8626560320469259



```
In [621... # Choice important features
important_features = ["ClientPeriod", "TotalSpent", "MonthlySpending", "HasInternetService_fiber optic"]
# Check score
rf_importances = RandomForestClassifier(max_depth=7, n_estimators=200, n_jobs=-1)
rf_importances.fit(x_train[important_features], y_train)
predict_rf_importances = rf_importances.predict_proba(x_test[important_features])[:, 1]

logreg = LogisticRegression()
logreg.fit(x_train[important_features], y_train)
predict_logreg_importances = logreg.predict_proba(x_test[important_features])[:, 1]

print(f"ROC-AUC RandomForest: {roc_auc_score(y_test, predict_rf_importances)}, ROC-AUC LogisticRegression: {roc_auc_score(y_test, predict_logreg_importances)}")
```

ROC-AUC RandomForest: 0.862405665438678, ROC-AUC LogisticRegression: 0.8449694195071354

```
In [622... # Remove duplicates for rows without non important features
# delete duplicates in data from train_file
# 1. delete full duplicates
x_data_important = pd.concat([x_data[important_features], y_data], axis=1)
x_data_important = x_data_important.drop_duplicates()

# 2. delete equivalent features with other targets!
mask_duplicate_rows = x_data_important.drop(target_colname, axis=1).duplicated()
x_data_important = x_data_important[~mask_duplicate_rows].reset_index()
```

```
# select y_data, x_data
y_data = x_data_important[target_colname]
x_data_important = x_data_important[important_features]

# testfile data select only important features
x_data_important_testfile = x_data_testfile[important_features]

# Train/Test split
x_train, x_test, y_train, y_test = train_test_split(x_data_important, y_data, test_size=0.2,
                                                    random_state=42,
                                                    x_data_important.shape, y_data.shape, x_train.shape, x_test.shape, y_train.shape, y_test.shape)
```

Out[622]: ((5230, 14), (5230,)), (4707, 14), (523, 14), (4707,)), (523,))

Models analysis

In [623...

```
# Selection of hyperparameter C
from sklearn.linear_model import LogisticRegressionCV
hparameter_C = [100, 10, 1, 0.1, 0.01, 0.001]
log_reg_cv = LogisticRegressionCV(Cs=hparameter_C, scoring='roc_auc', refit=True, cv=5)
log_reg_cv.fit(x_train, y_train)

# get score roc_auc:
auc_score_v1 = log_reg_cv.score(x_test, y_test)
predict_log_reg_cv = log_reg_cv.predict_proba(x_test)[:,:1]
auc_score_v2 = roc_auc_score(y_test, predict_log_reg_cv)

print(f"ROC AUC Score. calc score v1:{auc_score_v1}, calc score v2: {auc_score_v2}, v1==v2: True")
print(f"Best C parameters from ({hparameter_C}): {log_reg_cv.C_}")
```

ROC AUC Score. calc score v1:0.8602977667493796, calc score v2: 0.8602977667493796, v1==v2: True
Best C parameters from ([100, 10, 1, 0.1, 0.01, 0.001]): [1.]

In [624...

```
# Select the best parameters for RandomForestClassifier
from sklearn.model_selection import GridSearchCV
# set parameters
param_grid = {
    "max_depth": np.arange(3, 10, 3),
    "n_estimators": np.arange(200, 400, 50),
    "min_samples_split": np.arange(2, 6, 2),
    "bootstrap": [True, False],
}
# create and fit GridSearchCV model
rfc=RandomForestClassifier(n_estimators=200, random_state=42, criterion="entropy")
model = GridSearchCV(rfc, param_grid, cv=5, scoring='roc_auc', refit=True, n_jobs=-1)
model.fit(x_train, y_train)

# Show best parameters
print("Best ROC-AUC metric:", model.best_score_)
print("Best estimator's parameters:", model.best_params_)
```

Fitting 5 folds for each of 48 candidates, totalling 240 fits
Best ROC-AUC metric: 0.8435295790339079
Best estimator's parameters: {'bootstrap': True, 'max_depth': 6, 'min_samples_split': 2, 'n_estimators': 300}

Stacking

In [625...

```

catboost_model = catboost.CatBoostClassifier(n_estimators=200, silent=True, eval_me
rand_forest_3 = RandomForestClassifier(max_depth=3, n_estimators=200, bootstrap=Tru
rand_forest_7 = RandomForestClassifier(max_depth=7, n_estimators=200, bootstrap=Tru
rand_forest_10 = RandomForestClassifier(max_depth=10, n_estimators=200, bootstrap=1
# rand_forest_3 = RandomForestClassifier(max_depth=3, n_estimators=200, 'bootstrap
# rand_forest_7 = RandomForestClassifier(max_depth=7, n_estimators=200, n_jobs=-1)
# rand_forest_10 = RandomForestClassifier(max_depth=10, n_estimators=200, n_jobs=-1

models = [
    catboost_model,
    rand_forest_3,
    rand_forest_7,
    rand_forest_10,
]

train_features_model_predicts = np.zeros(x_data_important.shape[0])[...,None]
test_features_model_predicts = np.zeros(x_data_important_testfile.shape[0])[...,None]

for model in models:
    print(f"model: {model}")
    model.fit(x_data_important, y_data)
    cv_predict = cross_val_predict(model, x_data_important, y_data, cv=5, method="p

    predict_test = model.predict_proba(x_data_important_testfile)
    train_features_model_predicts = np.concatenate((train_features_model_predicts,
    test_features_model_predicts = np.concatenate((test_features_model_predicts, pr

train_features_model_predicts = train_features_model_predicts[:,1:]
test_features_model_predicts = test_features_model_predicts[:,1:]

stacking_model = LogisticRegression(penalty='l2', C=1, max_iter=30)
stacking_model.fit(train_features_model_predicts, y_data)

predict_stacking_model_proba = stacking_model.predict_proba(test_features_model_pre

# Write submission
predict_df = pd.DataFrame(predict_stacking_model_proba)
predict_df = pd.concat([pd.DataFrame(predict_df.index.values), predict_df], axis=1)
predict_df.to_csv(PATH_DATASET + "stacking_important_submission.csv", sep=",", inde

model: <catboost.core.CatBoostClassifier object at 0x00000297989E66C8>
model: RandomForestClassifier(criterion='entropy', max_depth=3, n_estimators=200,
                             n_jobs=-1)
model: RandomForestClassifier(criterion='entropy', max_depth=7, n_estimators=200,
                             n_jobs=-1)
model: RandomForestClassifier(criterion='entropy', max_depth=10, n_estimators=200,
                             n_jobs=-1)

```

YOUR RECENT SUBMISSION



stacking_important_submission.csv

Submitted by Черныш Иван · Submitted a few seconds ago

Score: 0.85174